Rubric

Personal morning paper
Each candidate presents their product to a judge, who marks it against this rubric. Eight criteria, 20 points each, plus bonus points.

Scoring

Score each criterion from 0 to 20. Find the band that best matches what you see, then pick a score inside it.

AI architecture

Score	What it looks like	Tick
17-20	Everything in 13-16, and the paper still comes out when a source or the AI fails.	☐
13-16	Clear steps, sources fetched in parallel, and each AI call has one job with a checked output. But if a source or the AI fails, no paper comes out.	☐
9-12	Clear steps from gathering to finished paper. But sources are fetched one after another, or one AI call does several jobs.	☐
5-8	Some separate steps, but they run one after another. AI output isn't checked. One failure stops the paper.	☐
0-4	One big prompt does everything. Steps run one after another. One failure stops the paper.	☐

Token optimisation

Score	What it looks like	Tick
17-20	Everything in 13-16, and knows the tokens and cost of one paper.	☐
13-16	AI only where it adds value, and fetched articles and summaries are reused across readers with similar interests. Only a rough guess of the cost.	☐
9-12	AI is used only where it adds value; simple work (sorting, removing duplicates) is done in code. Nothing is reused across readers. Only a rough guess of the cost.	☐
5-8	Some filtering happens before the AI, but every reader's paper is made from scratch. Can't say what a paper costs.	☐
0-4	The AI reads and rewrites everything it finds. The same calls are repeated. Can't say what a paper costs.	☐

Web search optimisation

Score	What it looks like	Tick
17-20	Everything in 13-16, and knows how many searches one paper needs and what they cost.	☐
13-16	A topic is searched once and the results are shared by every reader who follows it.	☐
9-12	Searches run in parallel and results are cached for the day, but searches aren't shared across readers.	☐
5-8	Searches run in parallel, but every reader gets their own searches, even for the same topic.	☐
0-4	Searches run one after another for each reader. The same search is repeated. Can't say how many searches a paper needs.	☐

Product quality

Score	What it looks like	Tick
17-20	Every feature works end to end. No story repeats or inventions. The paper arrives on time, once.	☐
13-16	Facts correct, sources referred to, no repeats. But a feature sometimes fails, or the paper arrives late or twice.	☐
9-12	Facts match their sources and sources are named, but stories repeat (in the same paper or day after day).	☐
5-8	Core features work, but some stories have wrong facts or don't refer their source.	☐
0-4	Features missing or broken. Stories have invented or wrong facts.	☐

Product experience

Score	What it looks like	Tick
17-20	Everything in 13-16, plus extra useful features. Proper formatting everywhere. Everything works and feels smooth.	☐
13-16	Sign-up and onboarding are easy. Good PDF and newsletter that look like a real paper. Readable layout and easy-to-find feedback. A good basic experience.	☐
9-12	Everything works, but getting to the paper is tricky and not intuitive. Can't find the feedback option or the "more of this, less of this" option.	☐
5-8	Only one part works properly. The PDF isn't formatted correctly.	☐
0-4	Nothing works properly: sign-up, onboarding, the paper, the newsletter and the PDF all have problems.	☐

Learning the reader's taste

Score	What it looks like	Tick
17-20	Everything in 13-16, and feedback form has great user experience to it.	☐
13-16	All feedback is reflected in the next newspaper.	☐
9-12	The next paper changes based on reader's feedback (like, skip, more or less of a topic) but some feedback was missed.	☐
5-8	Feedback is collected but doesn't change the next paper.	☐
0-4	No learning. The paper only uses the onboarding answers.	☐

Technical depth

Score	What it looks like	Tick
17-20	Explains every decision and its trade-offs, knows how the system behaves as readers grow, and finds their way around their own code with ease.	☐
13-16	Explains decisions, trade-offs, and how the system behaves as readers grow. Slow to find things in their own code.	☐
9-12	Explains the main decisions and their trade-offs, but not how the system behaves as readers grow.	☐
5-8	Can explain what the system does, but not why it was built that way.	☐
0-4	Can't explain how the system works or why it was built that way.	☐

Creativity

Score	What it looks like	Tick
17-20	Extra features and a great colour scheme. Feels complete and polished, with nothing to fault.	☐
13-16	A proper theme plus extra features beyond the brief that improve the paper.	☐
9-12	Basic features with a proper colour palette and a consistent theme.	☐
5-8	The basic features are implemented, and nothing more. Default look with no real design.	☐
0-4	Not even the basic features are implemented.	☐

Score sheet

Criterion	Score (out of 20)
AI architecture	
Token optimisation	
Web search optimisation	
Product quality	
Product experience	
Learning the reader's taste	
Technical depth	
Creativity	

Total (out of 160)	